

# structuring your workspace: DS & DE/MLE perspectives

Marco Morales

[marco.morales@columbia.edu](mailto:marco.morales@columbia.edu)

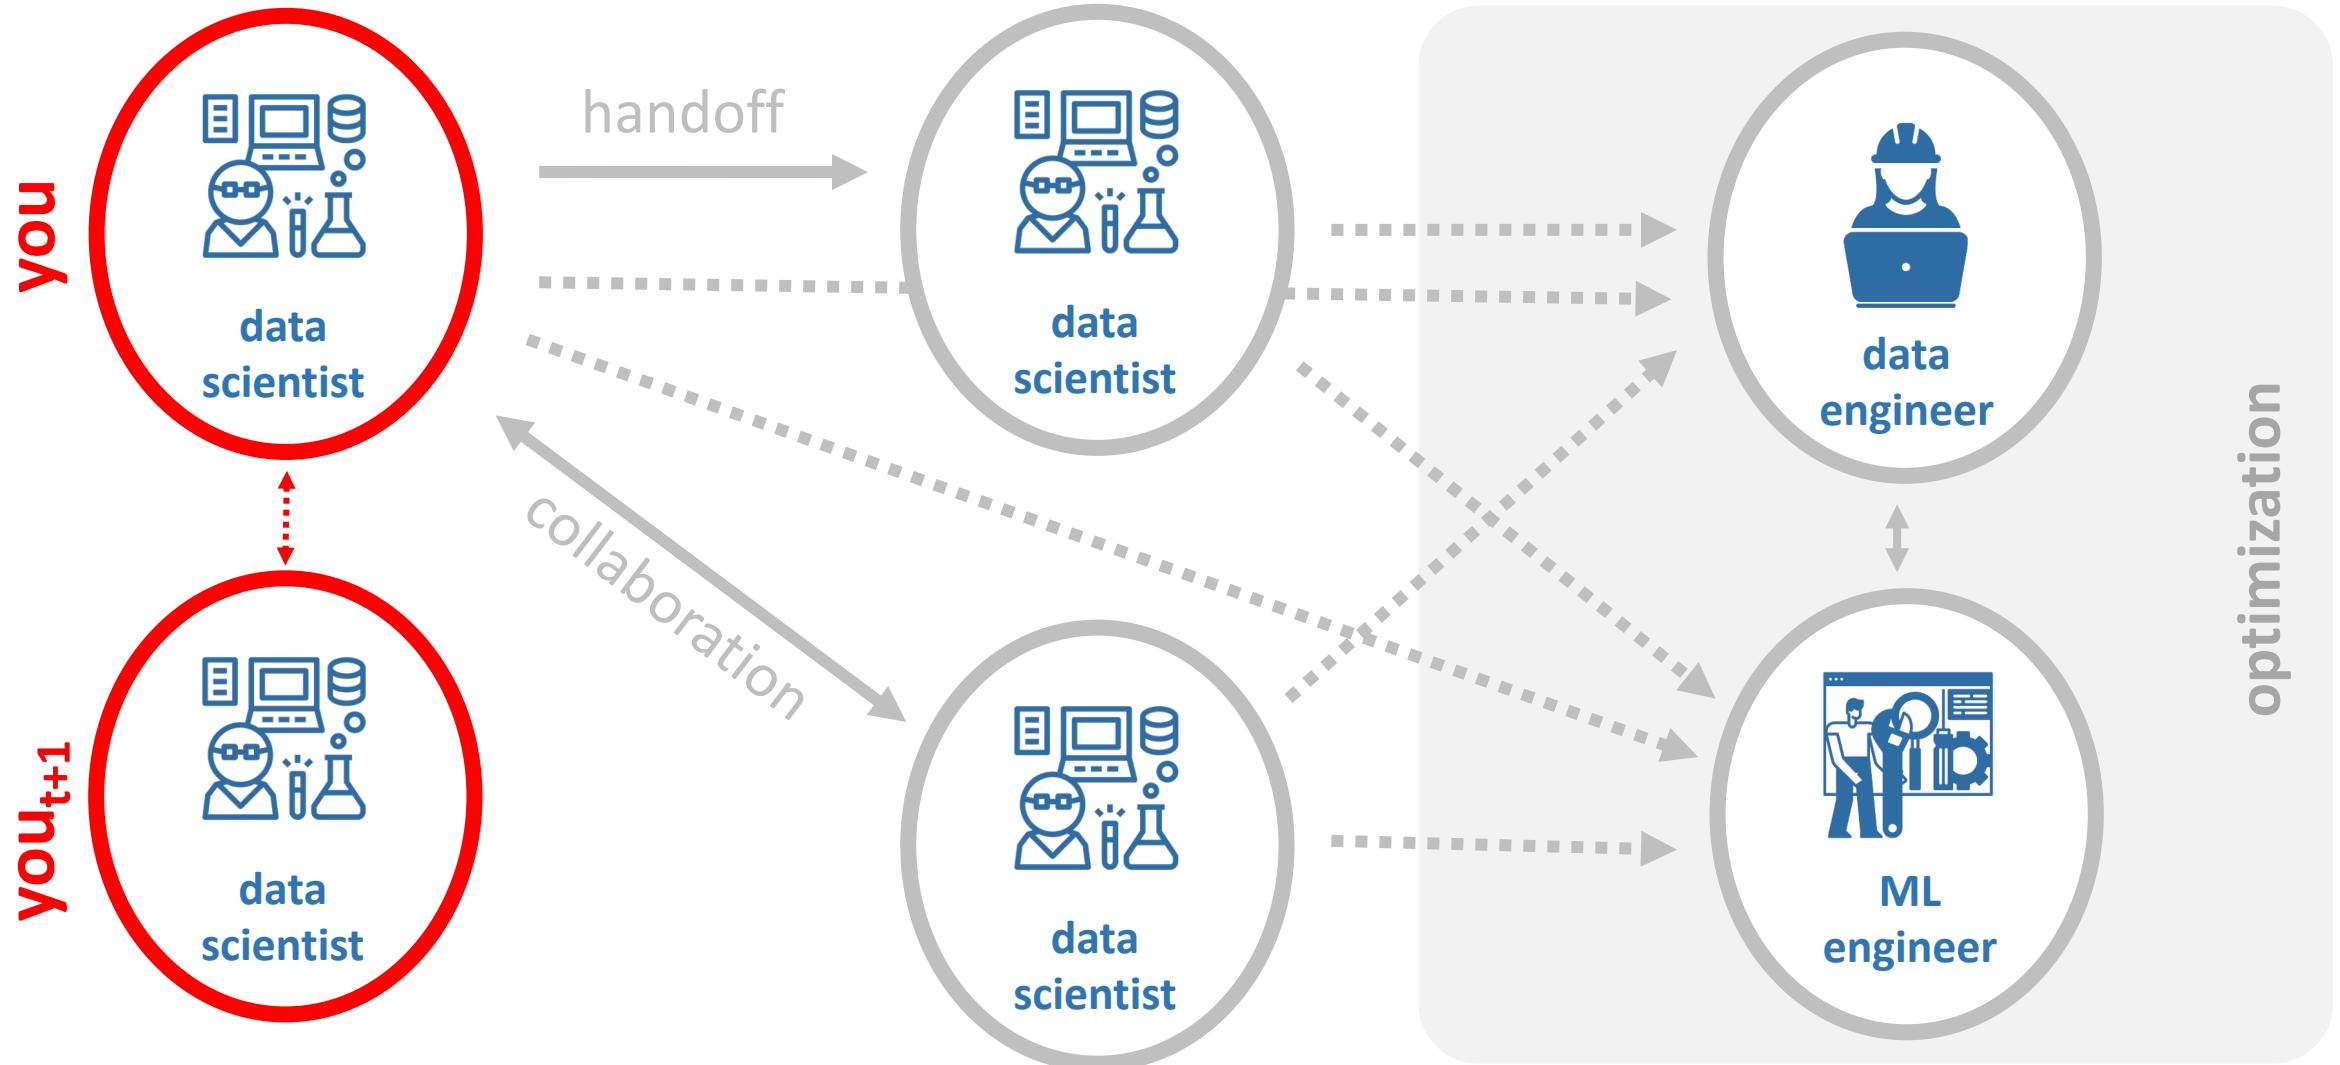
Nana Yaw Essuman

[ne2388@columbia.edu](mailto:ne2388@columbia.edu)

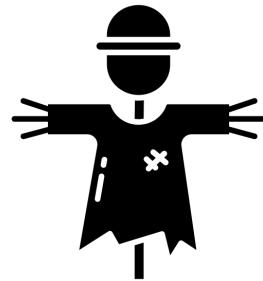
GR5069: Applied Data Science  
for Social Scientists

Spring 2026  
Columbia University

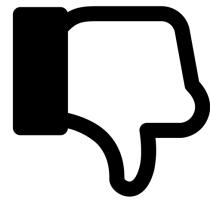
# workflow collaboration in Data Science



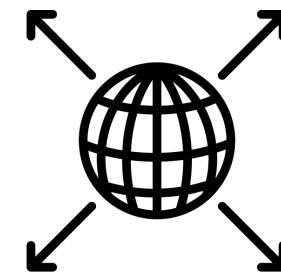
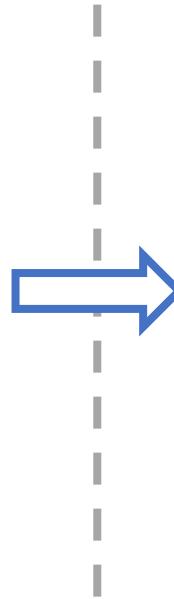
# recap: iteration to build Data Products



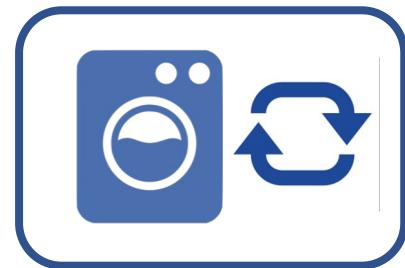
start small  
(MVP)



fail fast

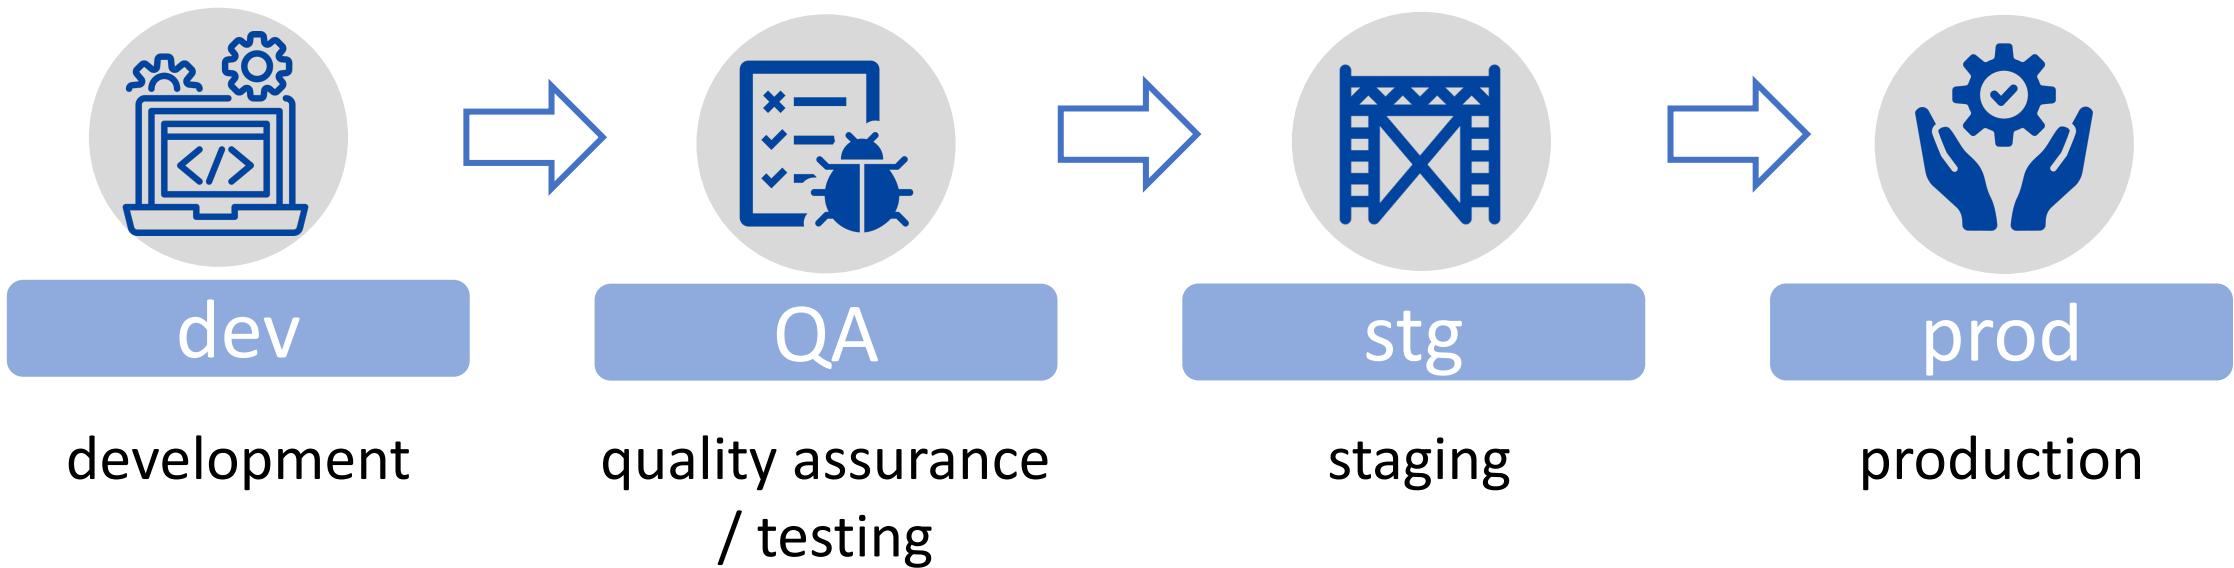


scale up

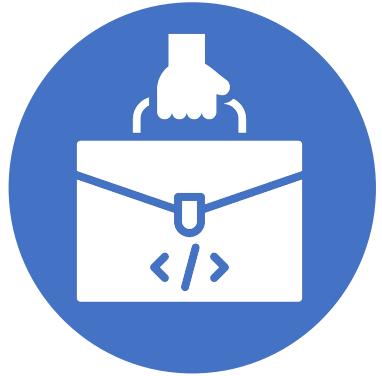


iterate

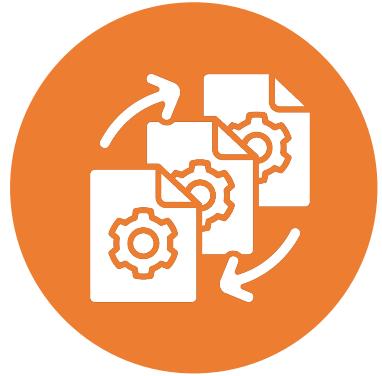
# working environments to build Data Products



# operational concepts in Data Science



**portability**



**replicability**



**scalability**

anyone should be able  
to **pick up where you  
left off** from any  
**machine**

anyone should be able  
to arrive at your **same  
results**

your prototype should  
also work for **larger  
data sets** and/or be on  
the path of **automation**

# operational concepts in Data Science



portability



replicability



scalability

what

- flexible references
- structured and documented code
- replicate original environment

- seamless handoff
- frictionless transitions across environments

- documentation: data, software, hardware, environments
- commented code
- no manual processes

- seamless examination, review or validation
- cordial troubleshooting
- harmonious optimization

- high quality code
  - flexible functions
  - modularized code
- 
- simplified review and validation
  - reduce time optimizing, automating and deploying

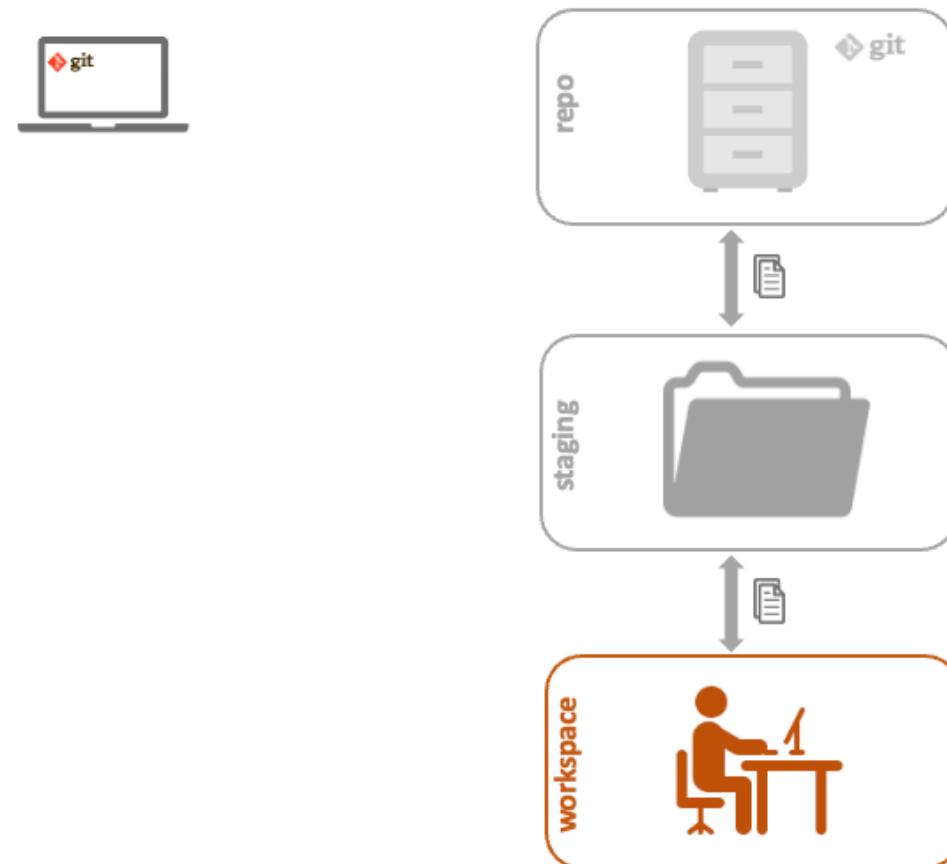
why

---

our focus :

---

# structuring your (local) workspace



# a thin layer to structure your (local) workspace

```
workspace/
|   -- src/           <- code
|   -- data/          <- inputs
|   -- reports/       <- outputs
|   -- references/    <- data dictionaries, explanatory materials
|
|   -- README.md
|   -- .gitignore
|   -- ToDo           <- bug fixes, future enhancements
|   -- LabNotebook    <- chronological records of project
```

a few principles to work by

# principle 1: code in a common location

```
workspace/
|
| -- src/
|   | -- data/                      <- code to read raw data
|   | -- features/                  <- code to transform data
|   | -- models/                   <- code to model data
|   | -- visualizations/          <- code to create visualizations
|   | -- functions/                <- scripts to centralize functions
|   | -- config/                  <- configuration files
|
| -- data/
|
| -- reports/
|
| -- references/
|
| -- README.md
| -- .gitignore
| -- ToDo
| -- LabNotebook
```

# principle 1: code in a common location

- use `src` to organize your **source code**
  - source code generates outputs; it's not generated
- use **one script per purpose**
- use **version control** to "update" your scripts
- use code to document "manual" changes
- call **additional scripts** as needed
- if too many functions, keep a **script with functions**

# principle 2: input raw data is immutable

```
workspace/
|   -- src/
|   -- data/
|       |   -- raw/           <- original, immutable data dump
|       |   -- external/      <- data from third party sources
|       |   -- interim/        <- intermediate transformed data (for validation)
|       |   -- processed/      <- final processed data
|
|   -- reports/
|
|   -- references/
|
|   -- README.md
|   -- .gitignore
|   -- ToDo
|   -- LabNotebook
```

## principle 2: input raw data is immutable

- **ALWAYS** keep your **raw data** as **immutable**
- keep **external data** separate and immutable
- (if/when needed) keep **interim data for validation**
- **processed data is ALWAYS replaceable!**
- all data should be linked to a script in **src**
- **document** origin of **raw & external data**

# principle 3: outputs are disposable

```
workspace/
|   -- src/
|   -- data/
|   -- reports/
|       |   -- notebooks/           <- documents synthesizing analysis, experiments, etc
|       |   -- figures/            <- images generated by code
|
|   -- references/
|
|   -- README.md
|   -- .gitignore
|   -- ToDo
|   -- LabNotebook
```

# principle 3: outputs are disposable

- use whichever document works best for your purpose:
  - Jupyter notebooks, R Markdown, scripts
- use notebooks to keep track of
  - **analyses**
  - **experiments**
  - **visualizations**
  - **tests**
  - **validation**
- **notebooks can be updated and are subject to change**

# principle 4: keep documentation

```
workspace/
|   -- src/
|   -- data/
|   -- reports/
|   -- references/           <- data dictionaries, explanatory materials
|   -- README.md
|   -- .gitignore
|   -- ToDo                  <- bug fixes, future enhancements
|   -- LabNotebook           <- chronological records of project
```

# principle 4: keep documentation

- **data**
  - **origin** and **schema** of your data sources (raw, first-to-third party)
  - data **dictionaries**
  - **other** relevant documentation
- **process** (LabNotebook)
  - what have you **tried** so far (and results)
  - what have you **not tried** but would if you had more time
- **ToDo**
  - **fixes** to current codebase
  - **improvements** to current codebase

# principle 5: keep records of your session

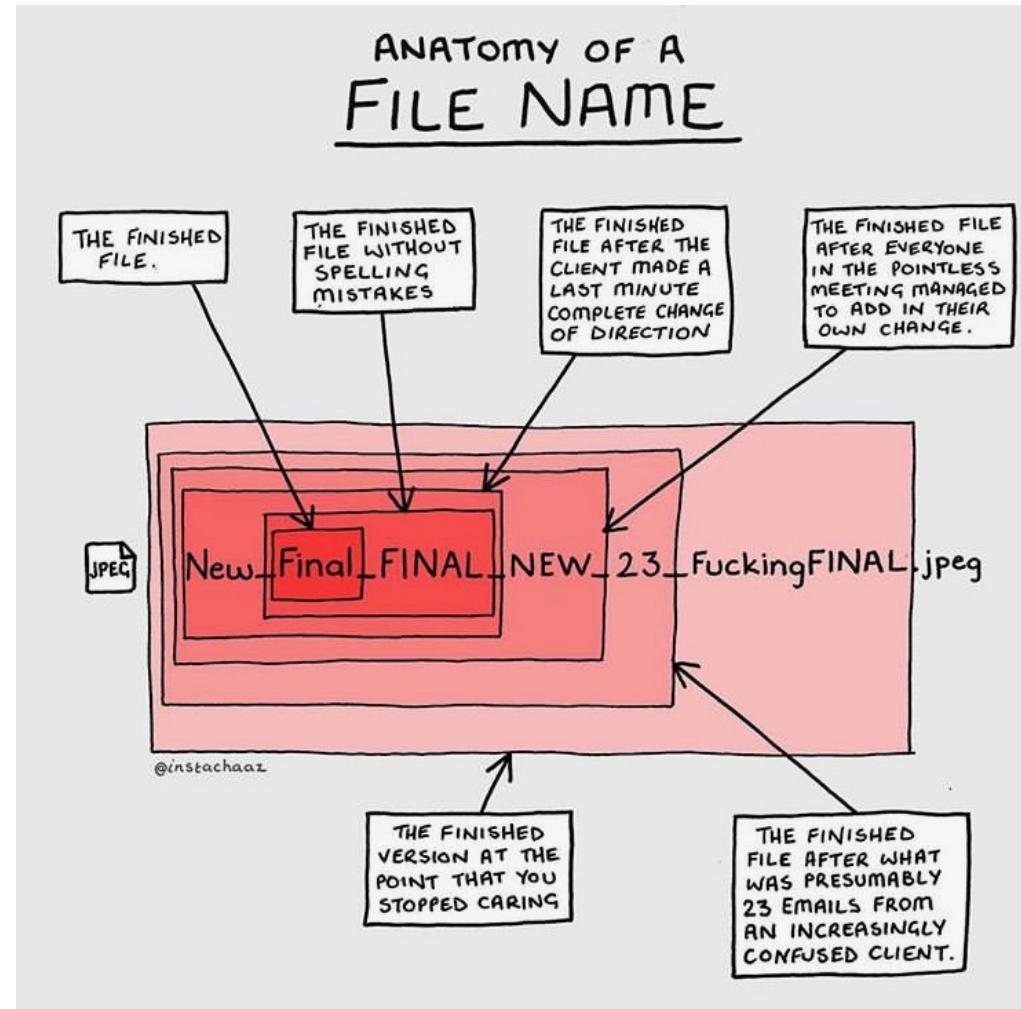
```
session_info.show(dependencies=True)

**Output:**
-----
matplotlib      3.7.2
numpy           1.24.3
pandas          2.0.3
session_info    1.0.0
-----
Python 3.11.5 | packaged by conda-forge | (main, Aug 27 2023, 03:34:09) [GCC 11.4.0]
Linux-5.15.0-1041-aws-x86_64-with-glibc2.35
Session information updated at 2024-02-14 15:30
-----
Dependencies
-----
asttokens       NA
certifi         2023.07.22
charset_normalizer 3.2.0
cycler          0.11.0
dateutil        2.8.2
```

# principle 5: keep records of your session

- always keep a record of your **session's**:
  - programming **language** and **version**
  - **libraries / packages** (versions, dependencies)
  - **OS**
  - **hardware**
- you'll need this info to **scale** or **hand off**:
  - pip freeze > requirements.txt

# principle 6: use informative script names



# principle 6: use informative script names

- provide detailed information about **purpose** of script / notebook in file name
  - use `clean_updated_data.py` instead of `untitled1.py`
- less is not more in this case
  - use `download_and_transform_dialy_data_updates.py` instead of `etl.py`
- use `_` as needed to **separate words**
  - use `download_and_transform_dialy_data_updates.py` instead of `donwloadandtransfrormdailydataupdates.py`
- **good file name:** informative + easy to read + appropriate length

# principle 7: always keep portability in mind

- use **relative paths** in your scripts "data\raw\file.csv" as opposed to "C:\users\data\raw\file.csv"
- or leverage your language's tools:

```
from pathlib import Path  
# Creating a path  
my_path = Path("data") / "raw" / "file.csv"
```

## principle 8: follow workplace *idioms*

- make sure to understand all *idioms* at your workplace
  - **best practices**
  - **conventions**
  - **styles**
  - **ad hoc rules**
- **specifics** may be unique to your workplace, general principles apply universally

# your (local) workspace in a nutshell...

```
workspace/
|
| -- src/
|   | -- data/
|   | -- features/
|   | -- models/
|   | -- visualizations/
|   | -- functions/
|   | -- config/
|
| -- data/
|   | -- raw/
|   | -- external/
|   | -- interim/
|   | -- processed/
|
| -- reports/
|   | -- notebooks/
|   | -- figures/
|
| -- references/
|
| -- README.md
| -- .gitignore
| -- ToDo
| -- LabNotebook
```

<- code  
<- code to read raw data  
<- code to transform data  
<- code to model data  
<- code to create visualizations  
<- scripts to centralize functions  
<- configuration files

<- inputs  
<- original, immutable data dump  
<- data from third party sources  
<- intermediate transformed data (for validation)  
<- final processed data

<- outputs  
<- documents synthesizing analysis, experiments, etc  
<- images generated by code

<- data dictionaries, explanatory materials

<- high-level project description/overview  
<- files to exclude from version control  
<- bug fixes, future enhancements  
<- chronological records of project

but not all workplaces are local...

# your workspace in real life

your  
workspace

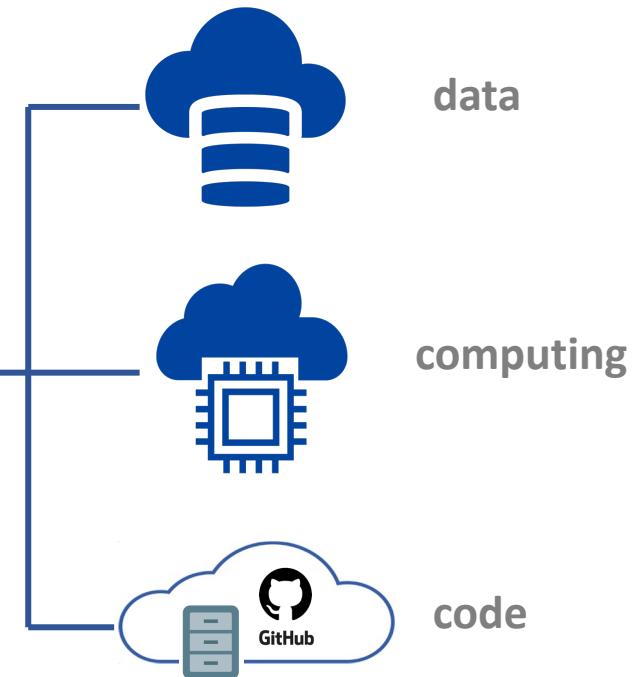
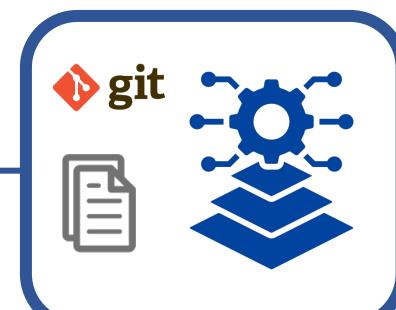


local

also  
your workspace



cloud-based



with a little help from our GitHub  
friends...

# what gets pushed to GitHub

```
workspace/
|   -- src/
|   |   -- data/           <- code to read raw data
|   |   -- features/       <- code to transform data
|   |   -- models/          <- code to model data
|   |   -- visualizations/ <- code to create visualizations
|   |   -- functions/        <- scripts to centralize functions
|   |   -- config/           <- configuration files
|
|   -- README.md
|   -- .gitignore
```

# what gets pushed to GitHub

- your `README.md` file should provide a **general overview** of your project
  - short (executive) **summary**
  - **structure** (navigation)
  - **instructions**
  - **limitations**
- yes, ALWAYS keep a meaningful `.gitignore` in your repo

# what gets pushed to GitHub

- data is **NEVER** pushed to GitHub!!!!!!
- {secret keys} are **NEVER** pushed to GitHub!!!!!!
- reports could live in GitHub (depends)
- references are transferred to GitHub **wiki**
- TODO is transferred to GitHub **issues / projects**

# why is a workspace structure needed?

- structure **removes friction** for
  1. any **first-time user** of your codebase – **collaborators, handoffs**
  2. the **you of the future** (if/when you return to a project)
- creates a **shared understanding** of where everything is – and where to look for anything!
- backbone of **portability, reproducibility, and scalability!**

# structuring your workspace: DS & DE/MLE perspectives

Marco Morales

[marco.morales@columbia.edu](mailto:marco.morales@columbia.edu)

Nana Yaw Essuman

[ne2388@columbia.edu](mailto:ne2388@columbia.edu)

GR5069: Applied Data Science  
for Social Scientists

Spring 2026  
Columbia University